

# Human Pose Estimation using Deep Learning

*A Project Report*

*Submitted to Jorhat Engineering College*

*in partial fulfillment of requirements for the award of degree*

*Bachelor of Technology*

*in*

*Computer Science and Engineering*

*by*

**Akangkshya Pathak(180712707001)**

**Apeksha Modi(180710007005)**

**Masoom Konwar(180710007034)**



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING  
JORHAT ENGINEERING COLLEGE JORHAT  
ASSAM  
August 2021**

## **CERTIFICATE FROM INSTITUTE**

This is to certify that the report entitled **Human Pose Estimation using Deep Learning** submitted by **Akangkshya Pathak** (180712707001), **Apeksha Modi** (180710007005) and **Masoom Konwar** (180710007005) under the guidance of Mr. Biswajit Sarmah in partial fulfillment of the requirements for the Degree of Computer Science And Engineering Final Year Project is a bona fide work carried out by them.

Signature of head of department:

**Dr. Rupam Baruah**  
Professor and Head  
Dept.of CSE  
Jorhat Engineering College  
Jorhat

## **CERTIFICATE FROM PROJECT GUIDE**

This is to certify that the report entitled **Human Pose Estimation using Deep Learning** submitted by **Akangkshya Pathak** (180712707001), **Apeksha Modi** (180710007005) and **Masoom Konwar** (180710007005) under the guidance of Mr. Biswajit Sarmah in partial fulfillment of the requirements for the Degree of Computer Science And Engineering Final Year Project is a bona fide work carried out by them.

Signature of Guide:

**Mr. Biswajit Sarmah**  
Asst. Professor  
Dept.of CSE  
Jorhat Engineering College  
Jorhat

## **DECLARATION**

We here by declared that the project work entitled **Human Pose Estimation using Deep Learning** ,submitted to the Jorhat Engineering College is an original work done with the guidance of Mr. Biswajit Sarmah, Assistant Professor, Dept. of Computer Science and Engineering, Jorhat Engineering College, and is submitted in partial fulfillment of the requirements for the Degree of B.E in Computer Science and Engineering. The result embodied in the report has not been submitted to any other University or Institute for the award of any degree or diploma.

Jorhat

**March 2022**

**Akangkshya Pathak**

**Apeksha Modi**

**Masoom Kownwar**

# Acknowledgement

We would like to express our heartfull gratitude to all the people who have help us from time to time during the course of this project without whom this work would never been accomplished. First and foremost we are very grateful to our HOD, **Dr. Rupam Baruah** Sir for including **Human Pose Estimation using Deep Learning** project in our course which greatly helps the students to enhance their practical development and coding skills and project guide **Mr. Biswajit Sarmah** Sir for his excellent idea and for having accepted our project. We are deeply indebted to both of them for their support, vision, advice, encouragement, passion and highly positive attitude for the project without whom our project would not been through. At last we shall remain indebted to our classmate and staffs of the Computer Science and Engineering Department for rendering their cooperation and everybody who have directly helped us in carrying out the project.

Akangkshya Pathak

Apeksha Modi

Masoom Konwar

# Abstract

Human Pose Estimation is a way of identifying and classifying the joints in the human body. It is a way to capture a set of coordinates for each joint for example head,neck, shoulder , hips, knees, ankles,hands etc which is known as a key point that can describe a pose of a person. The connection between these points is known as a pair and a skeleton is created by connecting these points in a hierarchical manner. In this project, we are trying to use Deep Learning, a method proposed by Alexander Toshev and Christian Szegedy published in the paper named “DeepPose: Human Pose Estimation via Deep Neural Networks” which uses DNN based regression to give coordinates of joints as output by taking images of humans doing some pose as an output. We tried to implement their model along with that we tried to make some of our own adjustment between the layers of the neural network to experiment with the outcome and accuracy of our goal is to improve the accuracy of Human pose estimation models and compare different CNN architectures with each other to see which one performs better in this model. While implementing the model, we faced issues with the pre existing datasets like LSP, FLIC for human pose estimation, hence we also designed a new data annotation tool which will help to prepare dataset specifically designed for this problem which may result improvement upon the overall accuracy of the various techniques. After which we implemented other CNN architectures, namely ResNet 34 and AlexNet to check the accuracy against the DeepPose CNN architecture and compared each model to find the best one.

# Contents

<b>Acknowledgement</b>	<b>i</b>
<b>Abstract</b>	<b>ii</b>
<b>List of Figures</b>	<b>vi</b>
<b>List of Tables</b>	<b>viii</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Pose Estimation problem . . . . .	1
1.2 Applications of Human Pose Estimation . . . . .	1
1.3 Dataset . . . . .	2
1.3.1 FLIC . . . . .	2
1.3.2 LSP . . . . .	2
1.4 Data Annotation Tool . . . . .	2
1.5 ResNet . . . . .	3
1.6 AlexNet . . . . .	3
<b>2 Literature Review</b>	<b>4</b>
2.1 Flowing ConvNets for Human Pose Estimation in Videos - 2015 . . .	4
2.1.1 Advantages . . . . .	4
2.1.2 Disadvantages . . . . .	5
2.2 ImageNet Classification with Deep Convolutional Neural Networks .	5
2.3 DeepPose: Human Pose Estimation via Deep Neural Networks - 2014	5
2.3.1 Advantages . . . . .	6
2.3.2 Disadvantages . . . . .	6

<b>3</b>	<b>Methodology</b>	<b>7</b>
3.1	Implementation of DeepPose Model . . . . .	7
3.1.1	Importing the libraries . . . . .	7
3.1.2	Reading the image . . . . .	9
3.1.3	Preprocessing . . . . .	9
3.1.4	Normalisation of coordinates . . . . .	10
3.1.5	Model Description . . . . .	10
3.1.6	Training model . . . . .	13
3.1.7	Evaluation . . . . .	14
3.1.8	Saving the models . . . . .	14
3.1.9	Constructing the final skeleton . . . . .	14
3.2	Data Annotation Tool . . . . .	15
3.2.1	introduction . . . . .	15
3.2.2	working . . . . .	15
3.3	Comparing with other CNN models . . . . .	16
3.3.1	introduction . . . . .	16
3.3.2	ResNet34 . . . . .	16
3.3.3	AlexNet . . . . .	22
<b>4</b>	<b>Results and Discussion</b>	<b>25</b>
4.1	Results . . . . .	25
4.2	Accuracy . . . . .	27
4.3	Loss . . . . .	34
4.4	Comparison with other CNN architectures . . . . .	41
4.4.1	Accuracy of the DeepPose model . . . . .	41
4.4.2	Accuracy of the ResNet34 model . . . . .	41
4.4.3	Accuracy of the AlexNet model . . . . .	42
<b>5</b>	<b>Future Objectives</b>	<b>43</b>
5.1	Build and real time motion capture application . . . . .	43
<b>6</b>	<b>Conclusion</b>	<b>44</b>





# List of Figures

- Fig 1 - Training Image Example
- Fig 2 - CNN model
- Fig 3 - DeepPose Model Description 1
- Fig 4 - DeepPose Model Description 2
- Fig 5 - Data Annotation Tool
- Fig 6 - ResNet34 Model Description
- Fig 7 - AlexNet Model Description
- Fig 8 - Result Image
- Fig 9 - Model accuracy of left hip
- Fig 10 - Model accuracy of left knee
- Fig 11 - Model accuracy of left ankle
- Fig 12 - Model accuracy of left shoulder
- Fig 13 - Model accuracy of left elbow
- Fig 14 - Model accuracy of left wrist
- Fig 15 - Model accuracy of right wrist
- Fig 16 - Model accuracy of right elbow
- Fig 17 - Model accuracy of right shoulder
- Fig 18 - Model accuracy of right knee
- Fig 19 - Model accuracy of right ankle
- Fig 20 - Model accuracy of right hip

- Fig 21 - Model accuracy of neck
- Fig 22 - Model accuracy of head
- Fig 23 - Model loss of left hip
- Fig 24 - Model loss of left knee
- Fig 25 - Model loss of left ankle
- Fig 26 - Model loss of left shoulder
- Fig 27 - Model loss of left elbow
- Fig 28 - Model loss of left wrist
- Fig 29 - Model loss of right wrist
- Fig 30 - Model loss of right elbow
- Fig 31 - Model loss of right shoulder
- Fig 32 - Model loss of right knee
- Fig 33 - Model loss of right ankle
- Fig 34 - Model loss of right hip
- Fig 35 - Model loss of neck
- Fig 36 - The final training and testing accuracy of all the joints for the DeepPose Model
- Fig 37 - The final training and testing accuracy of all the joints for the ResNet34 Model
- Fig 38 - The final training and testing accuracy of all the joints for the AlexNet Model

# List of Tables

- Table 1 - Training accuracy of different joints
- Table 2 - Testing accuracy of different joints
- Table 3 - Testing accuracy of different joints

# **Chapter 1**

## **Introduction**

### **1.1 Pose Estimation problem**

The problem of human pose estimation, defined as the problem of localization of human joints, has enjoyed substantial attention in the computer vision community. Some of the challenges of this problem – strong articulations, small and barely visible joints, occlusions and the need to capture the context. This is a very new topic in the field of computer vision although some of the papers previously have tried to solve the problem but the research is still ongoing and a lot of work is still yet to be done. Some of the methods that are used to solve this problem are part based models where they try to localize each part and then estimate the joints in the intersecting parts and DNN based methods like deep-pose and some other techniques that use heat-map of the joint area and optical flow to locate the joint.

### **1.2 Applications of Human Pose Estimation**

Solving the human pose estimation problem will lead to use of it in various fields like motion capture and animation where users will be able to capture "motion capture data" from their smartphone cameras and transfer that data to computer generated characters and also it will be applicable to create full body AR filters. It can also be used to control robotic or cybernetic arms and complete humanoid bodies using video feed of human as a controller.

## **1.3 Dataset**

After researching various sources and reading several papers on this topic we found out that there are primarily two datasets that are used in almost every model to benchmark the performance of the models. They are FLIC and LSP, both of which are available on public domain so we downloaded and tried to understand the contents of the data.

### **1.3.1 FLIC**

FLIC is an abbreviation for Frames Labeled In Cinema, This dataset contains 5003 stills or screenshots from various hollywood movies. These images contain single and multiple person and a mat file that comes with it contains coordinates of every visible joints of the every person present in the scene.

### **1.3.2 LSP**

This dataset contains 2000 pose annotated images of mostly sports people. The images have been scaled such that the most prominent person is roughly 150 pixels in length. Each image has been annotated with 14 joint locations. Left and right joints are consistently labelled from a person-centric viewpoint.

## **1.4 Data Annotation Tool**

In machine learning, data annotation is the process of labelling data to show the outcome you want your machine learning model to predict. We are marking - labelling, tagging, transcribing, or processing - a dataset with the features we want our machine learning system to learn to recognize. Once our model is deployed, we want it to recognize those features on its own and make a decision or take some action as a result. Annotated data reveals features that will train our algorithms to identify the same features in data that has not been annotated. Data annotation is used in supervised learning and hybrid, or semi-supervised, machine learning models that involve supervised learning.

## 1.5 ResNet

ResNet, short for Residual Network is a specific type of neural network that was introduced in 2015 by Kaiming He, Xiangyu Zhang, Shaoqing Ren and Jian Sun in their paper “Deep Residual Learning for Image Recognition”. Mostly in order to solve a complex problem, we stack some additional layers in the Deep Neural Networks which results in improved accuracy and performance. The intuition behind adding more layers is that these layers progressively learn more complex features. For example, in case of recognising images, the first layer may learn to detect edges, the second layer may learn to identify textures and similarly the third layer can learn to detect objects and so on. But it has been found that there is a maximum threshold for depth with the traditional Convolutional neural network model Residual networks have become quite popular for image recognition and classification tasks because of their ability to solve vanishing and exploding gradients when adding more layers to an already deep neural network.

## 1.6 AlexNet

AlexNet is the name of a convolutional neural network (CNN) architecture, designed by Alex Krizhevsky in collaboration with Ilya Sutskever and Geoffrey Hinton, who was Krizhevsky’s Ph.D. advisor. In AlexNet’s first layer, the convolution window shape is  $11 \times 11$ . Since most images in ImageNet are more than ten times higher and wider than the MNIST images, objects in ImageNet data tend to occupy more pixels. Consequently, a larger convolution window is needed to capture the object. The convolution window shape in the second layer is reduced to  $5 \times 5$ , followed by  $3 \times 3$ . In addition, after the first, second, and fifth convolutional layers, the network adds maximum pooling layers with a window shape of  $3 \times 3$  and a stride of 2. Moreover, AlexNet has ten times more convolution channels than LeNet. After the last convolutional layer there are two fully-connected layers with 4096 outputs. These two huge fully-connected layers produce model parameters of nearly 1 GB. AlexNet uses the ReLU activation function.

# Chapter 2

## Literature Review

### 2.1 Flowing ConvNets for Human Pose Estimation in Videos - 2015

**Authors : Tomas Pfister, James Charles and Andrew Zisserman**

In the technique proposed in this paper, a Convolution Network architecture is proposed that is able to benefit from temporal context by combining information across the multiple frames using optical flow. To this end a new network architecture is proposed which regresses a confidence heatmap of joint position predictions; incorporates optical flow at a mid-layer to align heatmap predictions from neighbouring frames; and includes a final parametric pooling layer which learns to combine the aligned heatmaps into a pooled confidence map.

#### 2.1.1 Advantages

1. This architecture outperforms neural networks that regress joint coordinates directly.
2. This paper improves the accuracy from previous methods.



### **2.1.2 Disadvantages**

1. The algorithm implemented in this paper cannot identify z coordinates in photos.

## **2.2 ImageNet Classification with Deep Convolutional Neural Networks**

**Authors: Alex Krizhevsky, Ilya Sutskever and Geoffrey E. Hinton**

In the technique proposed in this paper, we take an image as an input and use a 8 layer deep neural network to identify the location of joints in the body of the person depicted. Then we take each of these coordinates and draw a bounding box around each of them. After this we take each of these bounding boxes and apply another deep neural network in them to find the coordinate of these joints. The coordinate found is taken as  $y$  predicted and we train it with the actual value of  $y$  as we have got from the dataset using a neural network in order to improve its efficiency.

## **2.3 DeepPose: Human Pose Estimation via Deep Neural Networks - 2014**

**Authors : Alexander Toshev and Christian Szegedy**

In the technique proposed in this paper, we take an image as an input and use a 7 layer deep neural network to identify the location of joints in the body of the person depicted. Then we take each of these coordinates and draw a bounding box around each of them. After this we take each of these bounding boxes and apply another deep neural network in them to find the coordinate of these joints. The coordinate found is taken as  $y$  predicted and we train it with the actual value of  $y$  as we have got from the dataset using a neural network in order to improve its efficiency

### **2.3.1 Advantages**

1. The algorithm mentioned in the paper is easy to understand.
2. The algorithm mentioned in the paper is easy to implement.
3. This technique follows a simple approach.

### **2.3.2 Disadvantages**

1. The algorithm implemented in this paper cannot identify z coordinates in photos.
2. The accuracy of the model is low compared to other techniques.

# Chapter 3

## Methodology

The steps followed to create the model are:

### 3.1 Implementation of DeepPose Model

#### 3.1.1 Importing the libraries

For the implementation of this paper, we use the following libraries :

##### **TensorFlow**

TensorFlow is a free and open-source software library for machine learning and artificial intelligence. It can be used across a range of tasks but has a particular focus on training and inference of deep neural networks.

##### **Matplotlib**

Matplotlib is a plotting library for the Python programming language and its numerical mathematics extension NumPy. It provides an object-oriented API for embedding plots into applications using general-purpose GUI toolkits like Tkinter, wxPython, Qt, or GTK.

##### **Numpy**

NumPy is a library for the Python programming language, adding support for large, multi-dimensional arrays and matrices, along with a large collection of high-level

mathematical functions to operate on these arrays.

## **OpenCv**

OpenCV (Open Source Computer Vision Library) is an open source computer vision and machine learning software library. OpenCV was built to provide a common infrastructure for computer vision applications and to accelerate the use of machine perception in the commercial products.

## **OS**

The OS module in Python provides functions for interacting with the operating system. OS comes under Python's standard utility modules. This module provides a portable way of using operating system-dependent functionality. The `*os*` and `*os.path*` modules include many functions to interact with the file system.

## **Scipy**

SciPy is a scientific computation library that uses NumPy underneath. SciPy stands for Scientific Python. It provides more utility functions for optimization, stats and signal processing. Like NumPy, SciPy is open source so we can use it freely.

## **Scikit-learn**

Scikit-learn (formerly scikits.learn and also known as sklearn) is a free software machine learning library for the Python programming language. It features various classification, regression and clustering algorithms including support-vector machines, random forests, gradient boosting, k-means and DBSCAN, and is designed to interoperate with the Python numerical and scientific libraries NumPy and SciPy. Scikit-learn is a NumFOCUS fiscally sponsored project.

## **Pickle**

Python pickle module is used for serializing and de-serializing a Python object structure. Any object in Python can be pickled so that it can be saved on disk. What

pickle does is that it “serializes” the object first before writing it to file. Pickling is a way to convert a python object (list, dict, etc.)

### 3.1.2 Reading the image

We read the image using “imread” from openCV. We resize the image to a suitable size to optimize memory usage and orient the image to a correct way if it isn’t oriented properly.

### 3.1.3 Preprocessing

As preprocessing step we first resized the image to 100x100 pixels to provide a reasonable amount of parameters to the neural network. Then we found that after resizing the images the coordinates denoting the joints are no longer in the place where they should be. Hence, we used normalization so that the joints are corrected according to the new dimensions of the images and after that the data is ready to be fed into the CNN model.

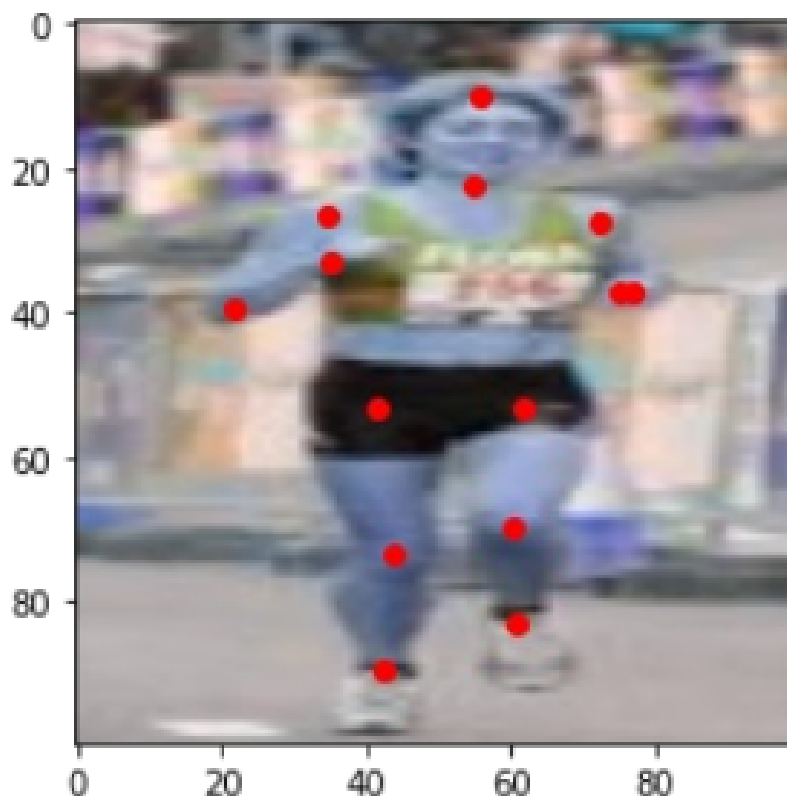


Fig 1

### 3.1.4 Normalisation of coordinates

We used lsp dataset to train the model. The dataset contained 2000 labeled images with a dictionary containing a matrix of 2000x14x3 dimension, where each image has 14 coordinates of the joints in that matrix but every image has their own size and shape and coordinates are also labeled according their shape. As we could not fit images with different shape to the CNN model, we resized and reshaped the image to 100x100 pixels and then used it as an input layer but because joint coordinates are labeled according to the shapes we had to normalize the coordinates of the joints to fit the new image shape so that the model does not predict wrong coordinates. We used the following formula:

$$coordinateX = \frac{coordinate_X}{dimension_X} * imagesize$$
$$coordinateY = \frac{coordinate_Y}{dimension_Y} * imagesize$$

Here,

$Coordinate_X = XCoordinate of the joint$

$Coordinate_Y = Ycoordinate of the joint$

$Dimension_X = Width of the image$

$Dimension_Y = Height of the image$

$Imagesize = Size of the input image$

### 3.1.5 Model Description

We implemented 14 separate CNN models to identify the different joints of the human body. The model contains 7 layers, denoted by C, a convoltional layer, Bn, a batch normalization layer, A, a relu activation layer, P, a max pooling layer and F, a fully connected layer. The layers C and F are the only layers containing learnable

parameters, while the rest of the layers are parameter free.

For C layers, the size is defined as:

size = width x height x depth

where,

the first two dimensions have a spatial meaning. The dimension named depth defines the number of filters. If we write the size of each layer in parenthesis, then the network can be described concisely as:

C ( 23 x 23 x 48 ) - Bn - A - P - C ( 11 x 11 x 128 ) - Bn - A - P - C ( 5 X 5 X 192 ) - A - C ( 5 X 5 X 128 ) - A - P - F ( 4096 ) - F ( 4096 )

The filter size for the first two C layers is 11 x 11 and for the remaining two layers, the filter size is 3 x 3. Same padding is used to preserve the size of all convolutional layers. Max pooling is applied after three layers and contributes to increased performance despite the reduction in resolution. The input to the net is an image of 100 x 100 x 3 dimensions which, via a stride of 4 is fed into the network. The output layer consists of two neurons. In this output layer, relu is used as an activation function and output from this layer is considered as final x and y coordinated of the model.




Fig 2

Model: "model\_7"

Layer (type)	Output Shape	Param #
input_8 (InputLayer)	[(None, 100, 100, 3)]	0
conv2d_35 (Conv2D)	(None, 23, 23, 48)	17472
batch_normalization_14 (BatchNormalization)	(None, 23, 23, 48)	192
tf.nn.relu_42 (TFOpLambda)	(None, 23, 23, 48)	0
max_pooling2d_21 (MaxPooling2D)	(None, 11, 11, 48)	0
conv2d_36 (Conv2D)	(None, 11, 11, 128)	743552
batch_normalization_15 (BatchNormalization)	(None, 11, 11, 128)	512
tf.nn.relu_43 (TFOpLambda)	(None, 11, 11, 128)	0
max_pooling2d_22 (MaxPooling2D)	(None, 5, 5, 128)	0
conv2d_37 (Conv2D)	(None, 5, 5, 192)	221376
tf.nn.relu_44 (TFOpLambda)	(None, 5, 5, 192)	0
conv2d_38 (Conv2D)	(None, 5, 5, 192)	331968
tf.nn.relu_45 (TFOpLambda)	(None, 5, 5, 192)	0
conv2d_39 (Conv2D)	(None, 5, 5, 128)	221312
tf.nn.relu_46 (TFOpLambda)	(None, 5, 5, 128)	0
max_pooling2d_23 (MaxPooling2D)	(None, 2, 2, 128)	0
flatten_7 (Flatten)	(None, 512)	0
dense_21 (Dense)	(None, 4096)	2101248
dense_22 (Dense)	(None, 4096)	16781312
dense_23 (Dense)	(None, 2)	8194
tf.nn.relu_47 (TFOpLambda)	(None, 2)	0

Fig 3



Total params: 20,427,138
Trainable params: 20,426,786
Non-trainable params: 352
None

Fig 4

### 3.1.6 Training model

For this process, we trained the model for 100 epochs for each joint using the lsp dataset different joints show different accuracy while training the list of final accuracy for each training set are as follows:

Joints	Accuracy
Right Knee	96.86%
Right ankle	98.33%
Right hip	97%
Left hip	98.14%
Left Knee	97.14%
Left ankle	99.14%
Right Wrist	97.86%
Right elbow	96.79%
Right shoulder	98.21%
Left shoulder	98.79%
Left elbow	97.57%
Left wrist	98.50%
Neck	99.29%
Head	99.64%

Table 1

### 3.1.7 Evaluation

The data set is tested by using 30% of the original LSP dataset, then we evaluated it for how well it performs in unseen data. The result are as follows.

Joints	Accuracy
Right Knee	86.83%
Right ankle	92.33%
Right hip	73.17%
Left hip	70.00%
Left Knee	80.50%
Left ankle	91.00%
Right Wrist	85.67%
Right elbow	85.50%
Right shoulder	86.21%
Left shoulder	88.83%
Left elbow	76.50%
Left wrist	73.17%
Neck	94.33%
Head	95.00%

Table 2

### 3.1.8 Saving the models

After training the model we saved the model using pickle module. For each joint we created a .sav file containing the weights and biases of the neurons of the model and it can be loaded at anytime in the future to predict the coordinates of the joints in any image containing a human figure.

### 3.1.9 Constructing the final skeleton

After we saved all the weights and biases for each joint we load all the weights and biases in different models for different joints. Then use those models to predict all the

joints and then connect the adjacent joints to construct the human skeleton to estimate the final pose of the human body.

## **3.2 Data Annotation Tool**

### **3.2.1 introduction**

In machine learning, data annotation is the process of labelling data to show the outcome you want your machine learning model to predict. We are marking - labelling, tagging, transcribing, or processing - a dataset with the features we want our machine learning system to learn to recognize. Once our model is deployed, we want it to recognize those features on its own and make a decision or take some action as a result.

### **3.2.2 working**

The dataset we used to train the model i.e lsp or leeds sports pose was very low resolution and poses were very extreme instead of being subtle so we decided to build a data annotation tool that can take images of people posing and annotate 14 joints in that pose with the help of an user it is a simple tool which uses web technologies such as html , css , javascript and node to load images from a back-end server to a front-end page and this page has 14 different markers which are color-coded according to the joints and these markers are all draggable these draggable markers can be dragged to the right position of the joints and when all the joints fall into the right positions we can save the coordinate and download a file containing the coordinates of the joints in the image. The tool writes in a .csv format file which then can be read using python as a numpy array. We then used this numpy array and the image in our previously built model to train/test using our previously built CNN models.



Fig 5

## 3.3 Comparing with other CNN models

### 3.3.1 introduction

After implementing the deep pose model with our own modifications and adjustments. We decided to replace the CNN model with some of the newer architectures available today like AlexNet and ResNet34 to see how well they perform in this problem and benchmark them with our own implementation

### 3.3.2 ResNet34

#### Implementing the ResNet 35 architecture

ResNet, short for Residual Network is a specific type of neural network that was introduced in 2015 by Kaiming He, Xiangyu Zhang, Shaoqing Ren and Jian Sun in their paper “Deep Residual Learning for Image Recognition”. Mostly in order

to solve a complex problem, we stack some additional layers in the Deep Neural Networks which results in improved accuracy and performance. The intuition behind adding more layers is that these layers progressively learn more complex features. For example, in case of recognising images, the first layer may learn to detect edges, the second layer may learn to identify textures and similarly the third layer can learn to detect objects and so on. But it has been found that there is a maximum threshold for depth with the traditional Convolutional neural network model. Residual networks have become quite popular for image recognition and classification tasks because of their ability to solve vanishing and exploding gradients when adding more layers to an already deep neural network. We replaced the model with ResNet 35 architectural constraints by changing the underlying layers and its parameters.

```
model = ResNet34()
print(model.summary())
```

Model: "ResNet34"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_5 (InputLayer)	[(None, 150, 150, 3)]	0	[]
zero_padding2d_4 (ZeroPadding2D)	(None, 156, 156, 3)	0	['input_5[0][0]']
conv2d_144 (Conv2D)	(None, 78, 78, 64)	9472	['zero_padding2d_4[0][0]']
batch_normalization_132 (Batch Normalization)	(None, 78, 78, 64)	256	['conv2d_144[0][0]']
activation_132 (Activation)	(None, 78, 78, 64)	0	['batch_normalization_132[0][0]']
max_pooling2d_4 (MaxPooling2D)	(None, 39, 39, 64)	0	['activation_132[0][0]']
conv2d_145 (Conv2D)	(None, 39, 39, 64)	36928	['max_pooling2d_4[0][0]']
batch_normalization_133 (Batch Normalization)	(None, 39, 39, 64)	256	['conv2d_145[0][0]']
activation_133 (Activation)	(None, 39, 39, 64)	0	['batch_normalization_133[0][0]']
conv2d_146 (Conv2D)	(None, 39, 39, 64)	36928	['activation_133[0][0]']
batch_normalization_134 (Batch Normalization)	(None, 39, 39, 64)	256	['conv2d_146[0][0]']
add_64 (Add)	(None, 39, 39, 64)	0	['batch_normalization_134[0][0]', 'max_pooling2d_4[0][0]']
activation_134 (Activation)	(None, 39, 39, 64)	0	['add_64[0][0]']

conv2d_147 (Conv2D)	(None, 39, 39, 64)	36928	['activation_134[0][0]']
batch_normalization_135 (Batch Normalization)	(None, 39, 39, 64)	256	['conv2d_147[0][0]']
activation_135 (Activation)	(None, 39, 39, 64)	0	['batch_normalization_135[0][0]']
conv2d_148 (Conv2D)	(None, 39, 39, 64)	36928	['activation_135[0][0]']
batch_normalization_136 (Batch Normalization)	(None, 39, 39, 64)	256	['conv2d_148[0][0]']
add_65 (Add)	(None, 39, 39, 64)	0	['batch_normalization_136[0][0]', 'activation_134[0][0]']
activation_136 (Activation)	(None, 39, 39, 64)	0	['add_65[0][0]']
conv2d_149 (Conv2D)	(None, 39, 39, 64)	36928	['activation_136[0][0]']
batch_normalization_137 (Batch Normalization)	(None, 39, 39, 64)	256	['conv2d_149[0][0]']
activation_137 (Activation)	(None, 39, 39, 64)	0	['batch_normalization_137[0][0]']
conv2d_150 (Conv2D)	(None, 39, 39, 64)	36928	['activation_137[0][0]']
batch_normalization_138 (Batch Normalization)	(None, 39, 39, 64)	256	['conv2d_150[0][0]']
add_66 (Add)	(None, 39, 39, 64)	0	['batch_normalization_138[0][0]', 'activation_136[0][0]']
activation_138 (Activation)	(None, 39, 39, 64)	0	['add_66[0][0]']
conv2d_151 (Conv2D)	(None, 20, 20, 128)	73856	['activation_138[0][0]']
batch_normalization_139 (Batch Normalization)	(None, 20, 20, 128)	512	['conv2d_151[0][0]']
activation_139 (Activation)	(None, 20, 20, 128)	0	['batch_normalization_139[0][0]']
conv2d_152 (Conv2D)	(None, 20, 20, 128)	147584	['activation_139[0][0]']
batch_normalization_140 (Batch Normalization)	(None, 20, 20, 128)	512	['conv2d_152[0][0]']
conv2d_153 (Conv2D)	(None, 20, 20, 128)	8320	['activation_138[0][0]']
add_67 (Add)	(None, 20, 20, 128)	0	['batch_normalization_140[0][0]', 'conv2d_153[0][0]']
activation_140 (Activation)	(None, 20, 20, 128)	0	['add_67[0][0]']
conv2d_154 (Conv2D)	(None, 20, 20, 128)	147584	['activation_140[0][0]']
batch_normalization_141 (Batch Normalization)	(None, 20, 20, 128)	512	['conv2d_154[0][0]']
activation_141 (Activation)	(None, 20, 20, 128)	0	['batch_normalization_141[0][0]']
conv2d_155 (Conv2D)	(None, 20, 20, 128)	147584	['activation_141[0][0]']
batch_normalization_142 (Batch Normalization)	(None, 20, 20, 128)	512	['conv2d_155[0][0]']
add_68 (Add)	(None, 20, 20, 128)	0	['batch_normalization_142[0][0]', 'activation_140[0][0]']
activation_142 (Activation)	(None, 20, 20, 128)	0	['add_68[0][0]']
conv2d_156 (Conv2D)	(None, 20, 20, 128)	147584	['activation_142[0][0]']

batch_normalization_143 (Batch Normalization)	(None, 20, 20, 128)	512	['conv2d_156[0][0]']
activation_143 (Activation)	(None, 20, 20, 128)	0	['batch_normalization_143[0][0]']
conv2d_157 (Conv2D)	(None, 20, 20, 128)	147584	['activation_143[0][0]']
batch_normalization_144 (Batch Normalization)	(None, 20, 20, 128)	512	['conv2d_157[0][0]']
add_69 (Add)	(None, 20, 20, 128)	0	['batch_normalization_144[0][0]', 'activation_142[0][0]']
activation_144 (Activation)	(None, 20, 20, 128)	0	['add_69[0][0]']
conv2d_158 (Conv2D)	(None, 20, 20, 128)	147584	['activation_144[0][0]']
batch_normalization_145 (Batch Normalization)	(None, 20, 20, 128)	512	['conv2d_158[0][0]']
activation_145 (Activation)	(None, 20, 20, 128)	0	['batch_normalization_145[0][0]']
conv2d_159 (Conv2D)	(None, 20, 20, 128)	147584	['activation_145[0][0]']
batch_normalization_146 (Batch Normalization)	(None, 20, 20, 128)	512	['conv2d_159[0][0]']
add_70 (Add)	(None, 20, 20, 128)	0	['batch_normalization_146[0][0]', 'activation_144[0][0]']
activation_146 (Activation)	(None, 20, 20, 128)	0	['add_70[0][0]']
conv2d_160 (Conv2D)	(None, 10, 10, 256)	295168	['activation_146[0][0]']
batch_normalization_147 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_160[0][0]']
activation_147 (Activation)	(None, 10, 10, 256)	0	['batch_normalization_147[0][0]']
conv2d_161 (Conv2D)	(None, 10, 10, 256)	590080	['activation_147[0][0]']
batch_normalization_148 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_161[0][0]']
conv2d_162 (Conv2D)	(None, 10, 10, 256)	33024	['activation_146[0][0]']
add_71 (Add)	(None, 10, 10, 256)	0	['batch_normalization_148[0][0]', 'conv2d_162[0][0]']
activation_148 (Activation)	(None, 10, 10, 256)	0	['add_71[0][0]']
conv2d_163 (Conv2D)	(None, 10, 10, 256)	590080	['activation_148[0][0]']
batch_normalization_149 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_163[0][0]']
activation_149 (Activation)	(None, 10, 10, 256)	0	['batch_normalization_149[0][0]']
conv2d_164 (Conv2D)	(None, 10, 10, 256)	590080	['activation_149[0][0]']
batch_normalization_150 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_164[0][0]']
add_72 (Add)	(None, 10, 10, 256)	0	['batch_normalization_150[0][0]', 'activation_148[0][0]']
activation_150 (Activation)	(None, 10, 10, 256)	0	['add_72[0][0]']
conv2d_165 (Conv2D)	(None, 10, 10, 256)	590080	['activation_150[0][0]']
batch_normalization_151 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_165[0][0]']

activation_151 (Activation)	(None, 10, 10, 256)	0	['batch_normalization_151[0][0]']
conv2d_166 (Conv2D)	(None, 10, 10, 256)	590080	['activation_151[0][0]']
batch_normalization_152 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_166[0][0]']
add_73 (Add)	(None, 10, 10, 256)	0	['batch_normalization_152[0][0]', 'activation_150[0][0]']
activation_152 (Activation)	(None, 10, 10, 256)	0	['add_73[0][0]']
conv2d_167 (Conv2D)	(None, 10, 10, 256)	590080	['activation_152[0][0]']
batch_normalization_153 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_167[0][0]']
activation_153 (Activation)	(None, 10, 10, 256)	0	['batch_normalization_153[0][0]']
conv2d_168 (Conv2D)	(None, 10, 10, 256)	590080	['activation_153[0][0]']
batch_normalization_154 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_168[0][0]']
add_74 (Add)	(None, 10, 10, 256)	0	['batch_normalization_154[0][0]', 'activation_152[0][0]']
activation_154 (Activation)	(None, 10, 10, 256)	0	['add_74[0][0]']
conv2d_169 (Conv2D)	(None, 10, 10, 256)	590080	['activation_154[0][0]']
batch_normalization_155 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_169[0][0]']
activation_155 (Activation)	(None, 10, 10, 256)	0	['batch_normalization_155[0][0]']



conv2d_171 (Conv2D)	(None, 10, 10, 256)	590080	['activation_156[0][0]']
batch_normalization_157 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_171[0][0]']
activation_157 (Activation)	(None, 10, 10, 256)	0	['batch_normalization_157[0][0]']
conv2d_172 (Conv2D)	(None, 10, 10, 256)	590080	['activation_157[0][0]']
batch_normalization_158 (Batch Normalization)	(None, 10, 10, 256)	1024	['conv2d_172[0][0]']
add_76 (Add)	(None, 10, 10, 256)	0	['batch_normalization_158[0][0]', 'activation_156[0][0]']
activation_158 (Activation)	(None, 10, 10, 256)	0	['add_76[0][0]']
conv2d_173 (Conv2D)	(None, 5, 5, 512)	1180160	['activation_158[0][0]']
batch_normalization_159 (Batch Normalization)	(None, 5, 5, 512)	2048	['conv2d_173[0][0]']
activation_159 (Activation)	(None, 5, 5, 512)	0	['batch_normalization_159[0][0]']
conv2d_174 (Conv2D)	(None, 5, 5, 512)	2359808	['activation_159[0][0]']
batch_normalization_160 (Batch Normalization)	(None, 5, 5, 512)	2048	['conv2d_174[0][0]']
conv2d_175 (Conv2D)	(None, 5, 5, 512)	131584	['activation_158[0][0]']
add_77 (Add)	(None, 5, 5, 512)	0	['batch_normalization_160[0][0]', 'conv2d_175[0][0]']
activation_160 (Activation)	(None, 5, 5, 512)	0	['add_77[0][0]']
conv2d_176 (Conv2D)	(None, 5, 5, 512)	2359808	['activation_160[0][0]']
batch_normalization_161 (Batch Normalization)	(None, 5, 5, 512)	2048	['conv2d_176[0][0]']

activation_161 (Activation)	(None, 5, 5, 512)	0	['batch_normalization_161[0][0]']
conv2d_177 (Conv2D)	(None, 5, 5, 512)	2359808	['activation_161[0][0]']
batch_normalization_162 (Batch Normalization)	(None, 5, 5, 512)	2048	['conv2d_177[0][0]']
add_78 (Add)	(None, 5, 5, 512)	0	['batch_normalization_162[0][0]', 'activation_160[0][0]']
activation_162 (Activation)	(None, 5, 5, 512)	0	['add_78[0][0]']
conv2d_178 (Conv2D)	(None, 5, 5, 512)	2359808	['activation_162[0][0]']
batch_normalization_163 (Batch Normalization)	(None, 5, 5, 512)	2048	['conv2d_178[0][0]']
activation_163 (Activation)	(None, 5, 5, 512)	0	['batch_normalization_163[0][0]']
conv2d_179 (Conv2D)	(None, 5, 5, 512)	2359808	['activation_163[0][0]']
batch_normalization_164 (Batch Normalization)	(None, 5, 5, 512)	2048	['conv2d_179[0][0]']
add_79 (Add)	(None, 5, 5, 512)	0	['batch_normalization_164[0][0]', 'activation_162[0][0]']
activation_164 (Activation)	(None, 5, 5, 512)	0	['add_79[0][0]']
average_pooling2d_4 (AveragePooling2D)	(None, 3, 3, 512)	0	['activation_164[0][0]']
flatten_4 (Flatten)	(None, 4608)	0	['average_pooling2d_4[0][0]']
dense_8 (Dense)	(None, 512)	2359808	['flatten_4[0][0]']
dense_9 (Dense)	(None, 2)	1026	['dense_8[0][0]']
=====			
Total params: 23,667,458			
Trainable params: 23,652,226			
Non-trainable params: 15,232			
=====			
None			

Fig 6

### 3.3.3 AlexNet

#### Implementing the AlexNet architecture

AlexNet is the name of a convolutional neural network (CNN) architecture, designed by Alex Krizhevsky in collaboration with Ilya Sutskever and Geoffrey Hinton, who was Krizhevsky's Ph.D. advisor. In AlexNet's first layer, the convolution window shape is  $11 \times 11$ . Since most images in ImageNet are more than ten times higher and wider than the MNIST images, objects in ImageNet data tend to occupy more pixels. Consequently, a larger convolution window is needed to capture the object. The convolution window shape in the second layer is reduced to  $5 \times 5$ , followed by

3×3. In addition, after the first, second, and fifth convolutional layers, the network adds maximum pooling layers with a window shape of 3×3 and a stride of 2. Moreover, AlexNet has ten times more convolution channels than LeNet. After the last convolutional layer there are two fully-connected layers with 4096 outputs. These two huge fully-connected layers produce model parameters of nearly 1 GB. AlexNet uses the ReLU activation function.

We replaced the model with AlexNet architectural constraints by changing the underlying layers and its parameters.

```
print(AlexNet.summary())
```

Model: "sequential\_14"

Layer (type)	Output Shape	Param #
conv2d_70 (Conv2D)	(None, 35, 35, 96)	34944
batch_normalization_70 (Batch Normalization)	(None, 35, 35, 96)	384
max_pooling2d_42 (MaxPooling2D)	(None, 17, 17, 96)	0
conv2d_71 (Conv2D)	(None, 17, 17, 256)	614656
batch_normalization_71 (Batch Normalization)	(None, 17, 17, 256)	1024
max_pooling2d_43 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_72 (Conv2D)	(None, 8, 8, 384)	885120
batch_normalization_72 (Batch Normalization)	(None, 8, 8, 384)	1536
conv2d_73 (Conv2D)	(None, 8, 8, 384)	1327488
batch_normalization_73 (Batch Normalization)	(None, 8, 8, 384)	1536
conv2d_74 (Conv2D)	(None, 8, 8, 256)	884992
batch_normalization_74 (Batch Normalization)	(None, 8, 8, 256)	1024
max_pooling2d_44 (MaxPooling2D)	(None, 3, 3, 256)	0

flatten_14 (Flatten)	(None, 2304)	0
dense_42 (Dense)	(None, 4096)	9441280
dropout_28 (Dropout)	(None, 4096)	0
dense_43 (Dense)	(None, 4096)	16781312
dropout_29 (Dropout)	(None, 4096)	0
dense_44 (Dense)	(None, 2)	8194
=====		
Total params: 29,983,490		
Trainable params: 29,980,738		
Non-trainable params: 2,752		
None		

Fig 7

# Chapter 4

## Results and Discussion

### 4.1 Results

After running our model we, get the final result as shown below.

Joints	Accuracy
Right Knee	86.83%
Right ankle	92.33%
Right hip	73.17%
Left hip	70.00%
Left Knee	80.50%
Left ankle	91.00%
Right Wrist	85.67%
Right elbow	85.50%
Right shoulder	86.21%
Left shoulder	88.83%
Left elbow	76.50%
Left wrist	73.17%
Neck	94.33%
Head	95.00%

Table 3

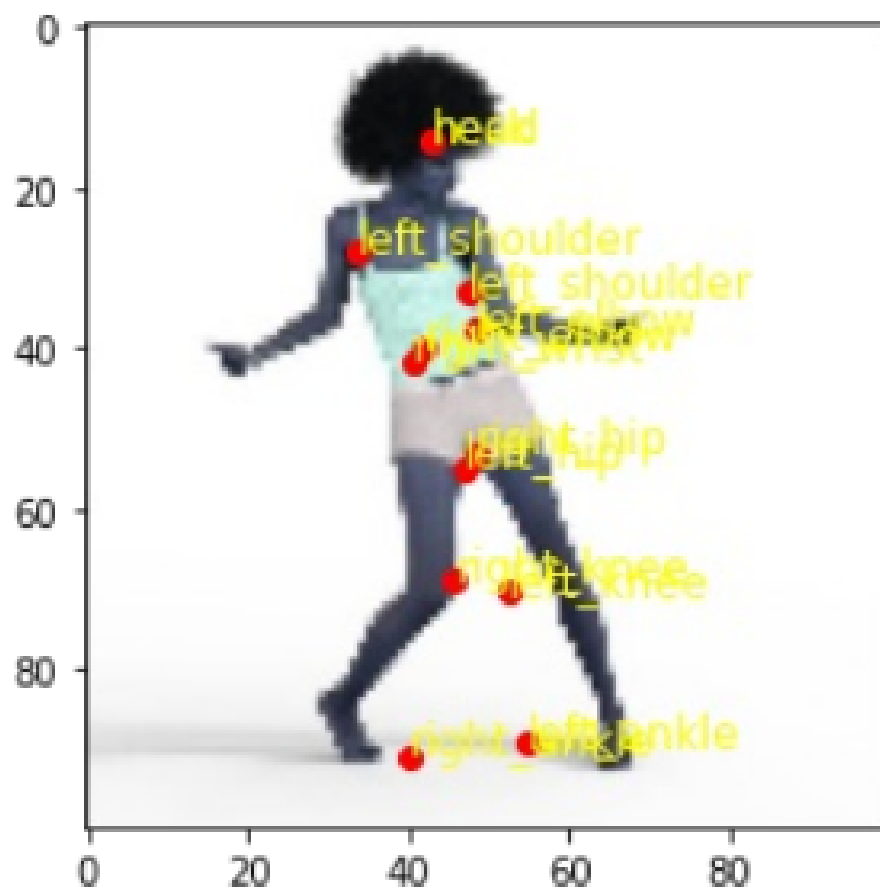


Fig 8

## 4.2 Accuracy

The accuracy is as follows:

### Left hip

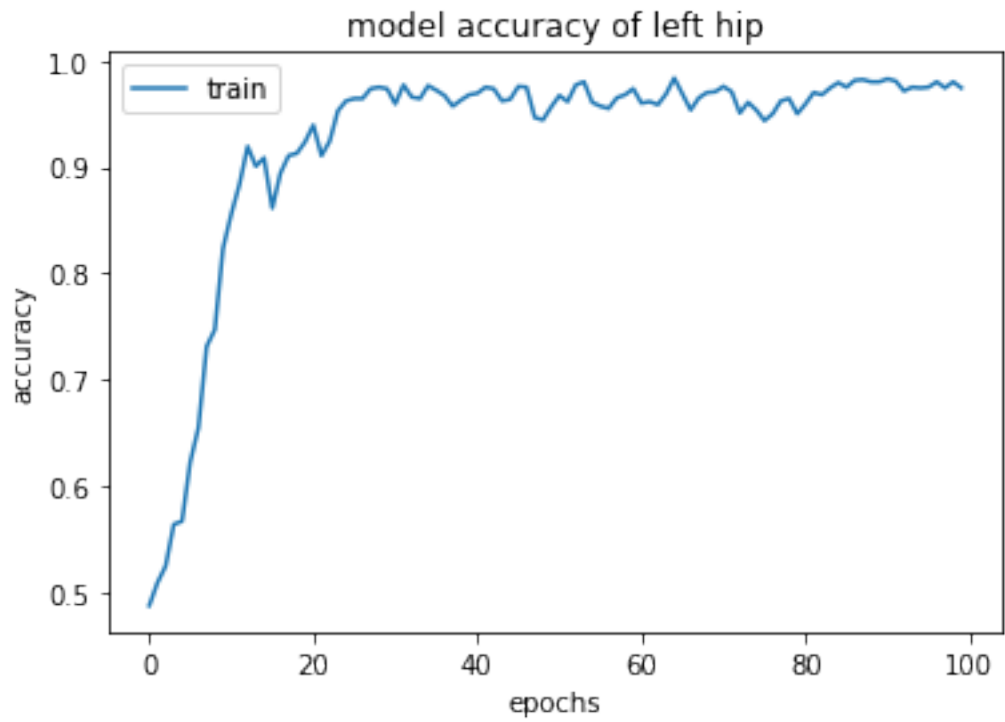


Fig 9

### Left Knee

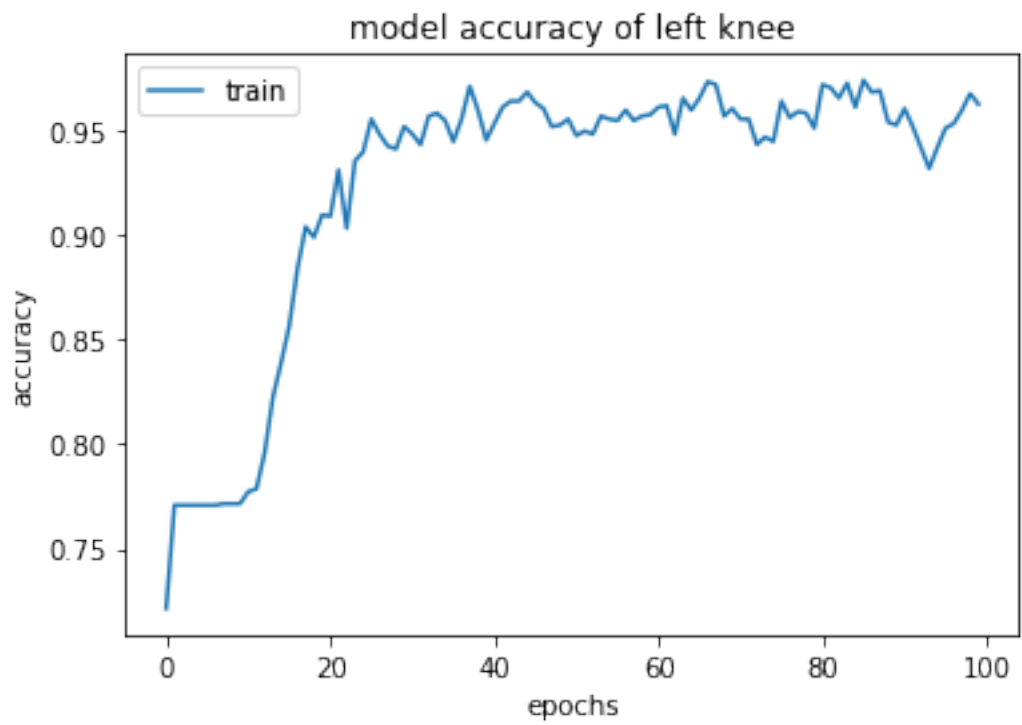


Fig 10

**Left ankle**

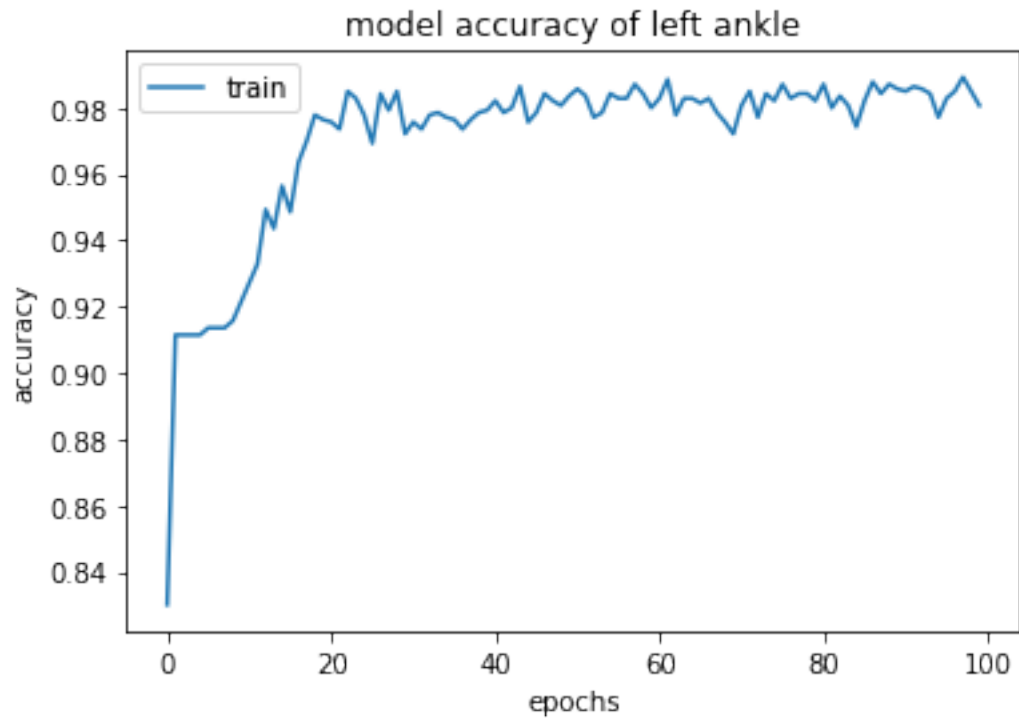


Fig 11

**Left shoulder**

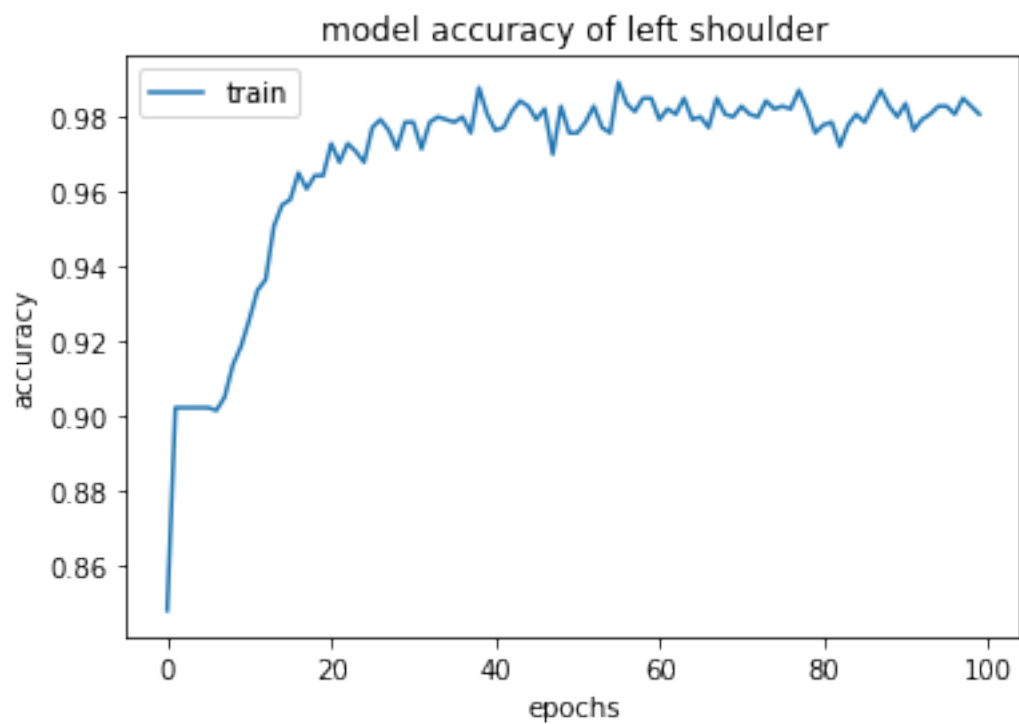


Fig 12



### Left elbow

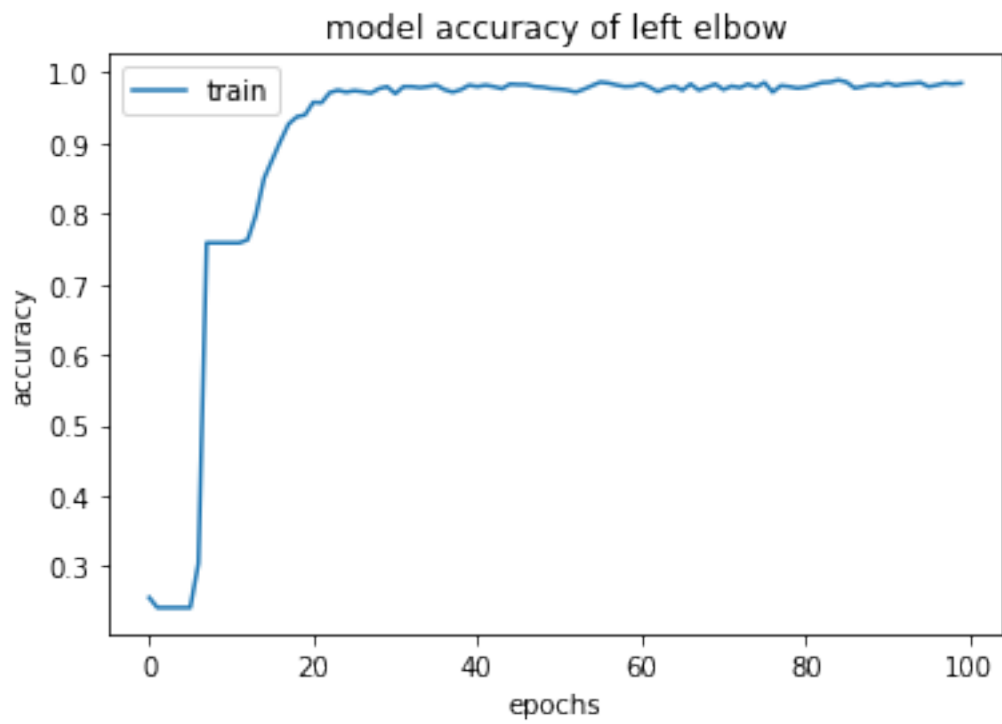


Fig 13

### Left wrist

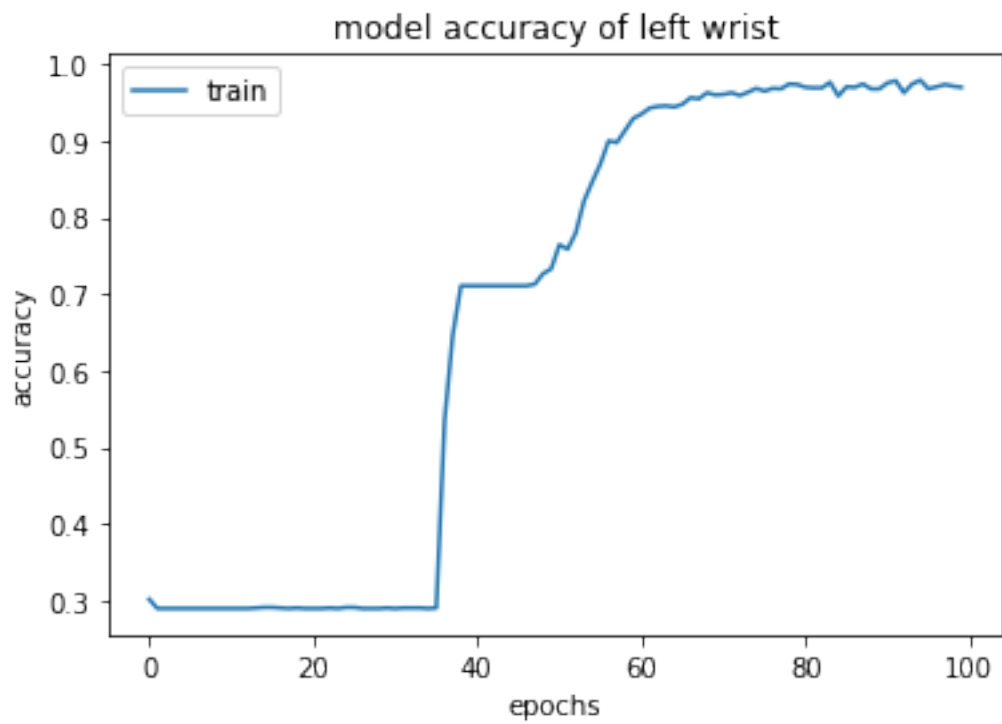


Fig 14

### Right Wrist

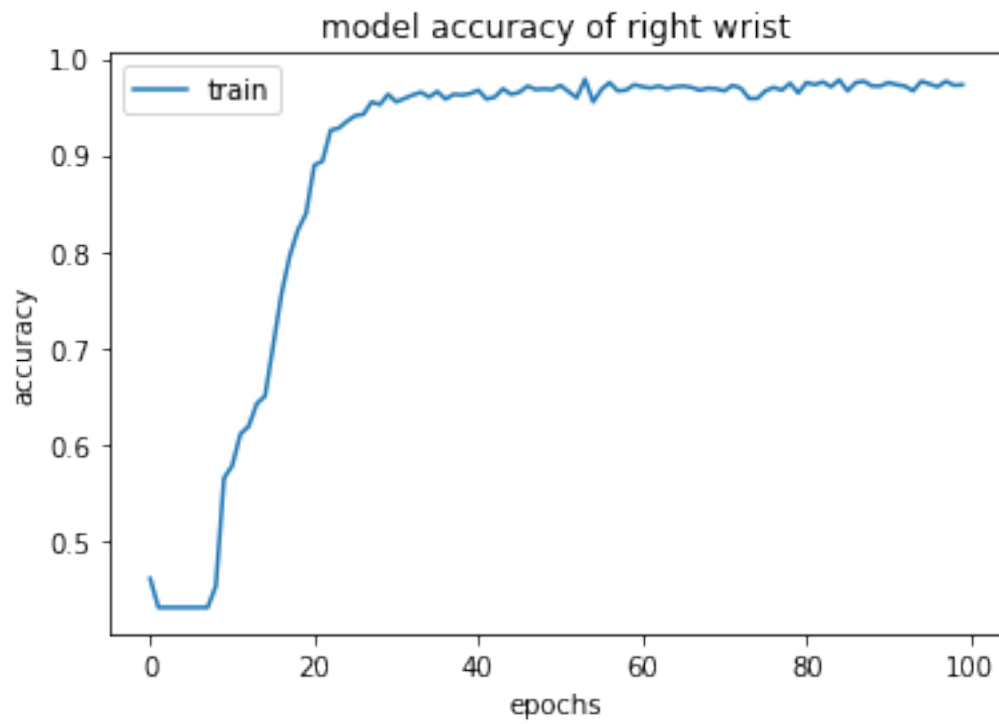


Fig 15

### Right elbow




Fig 16

### Right shoulder

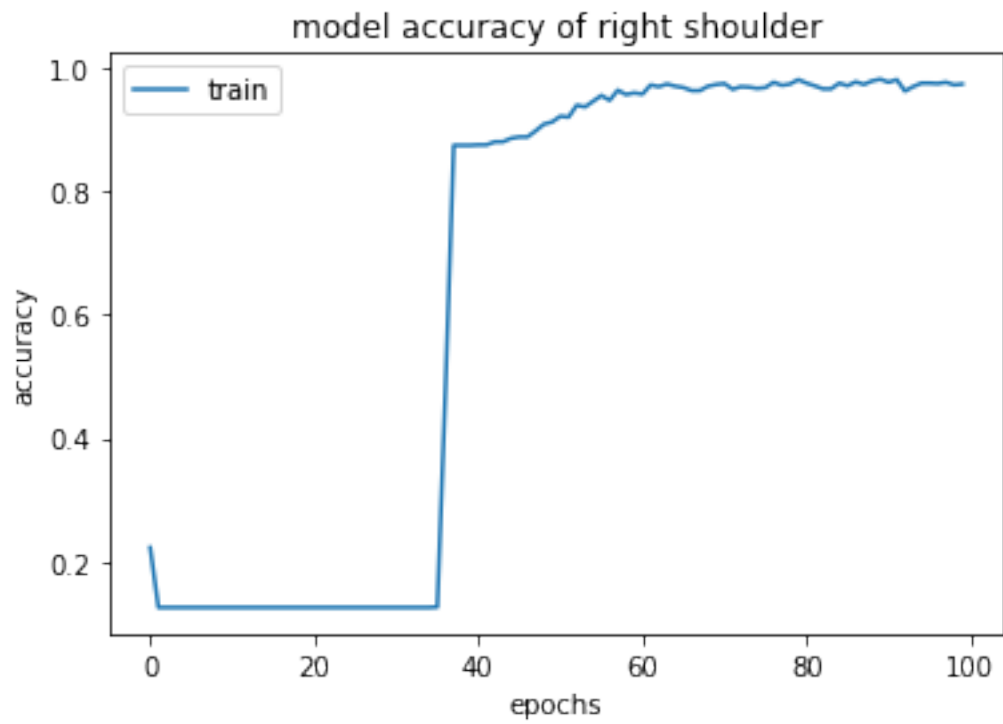


Fig 17

### Right Knee

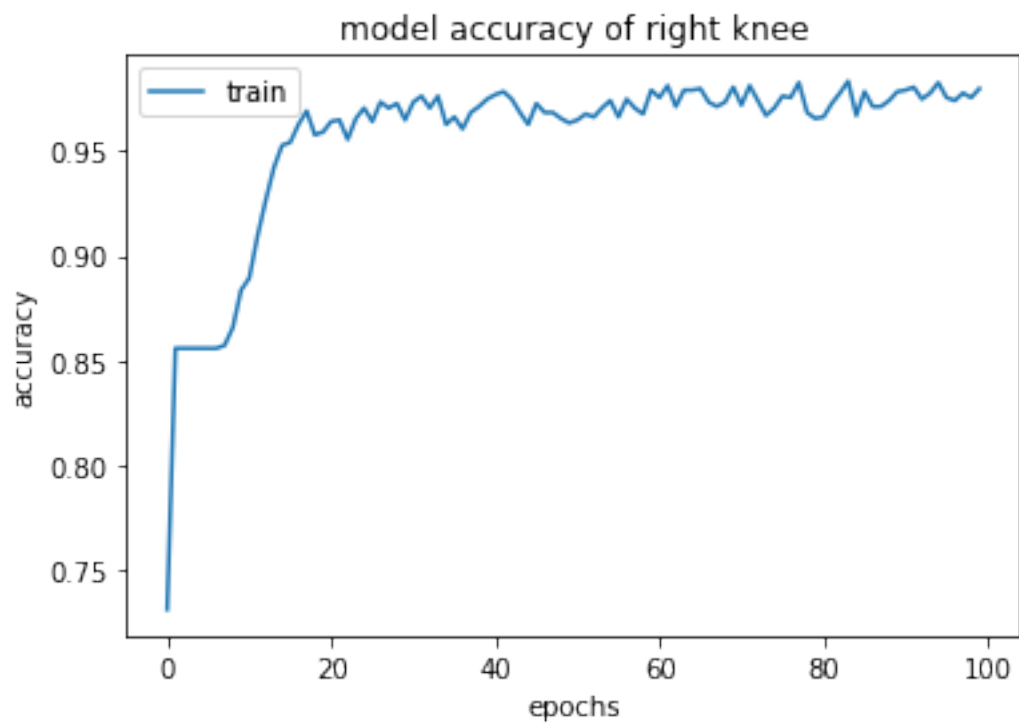


Fig 18

### Right ankle

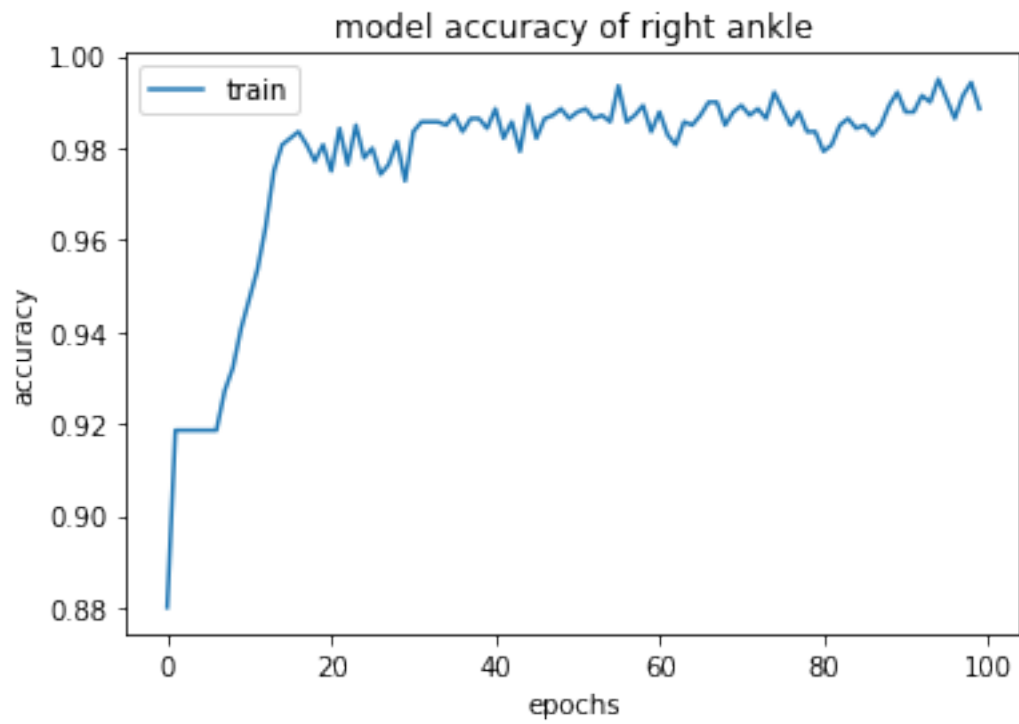


Fig 19

### Right hip

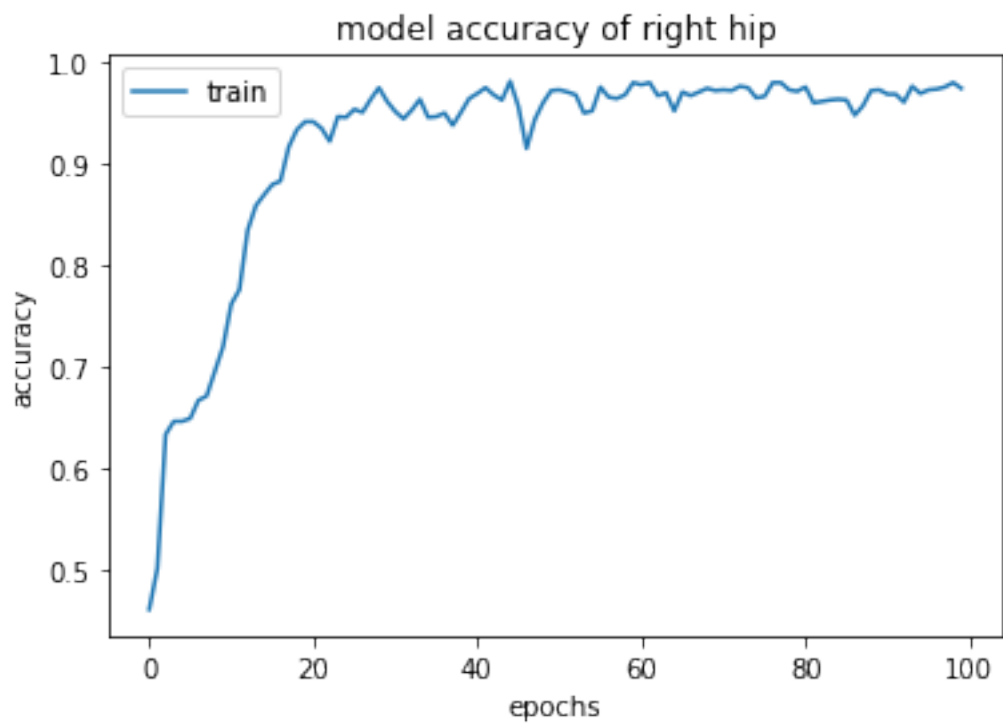


Fig 20

## Neck

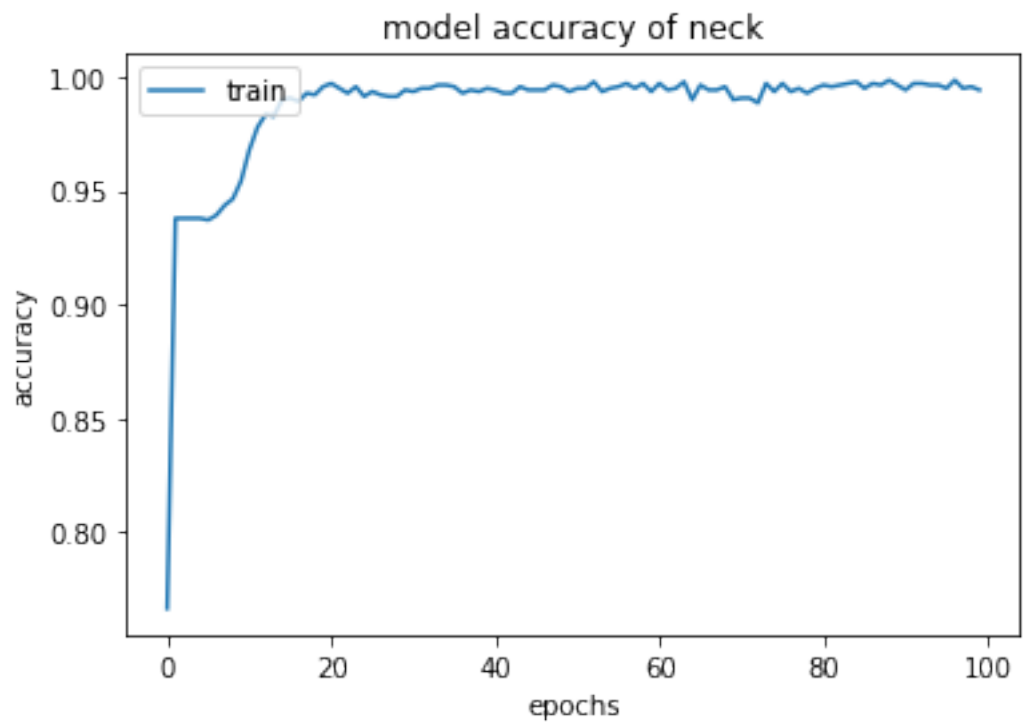


Fig 21

## Head

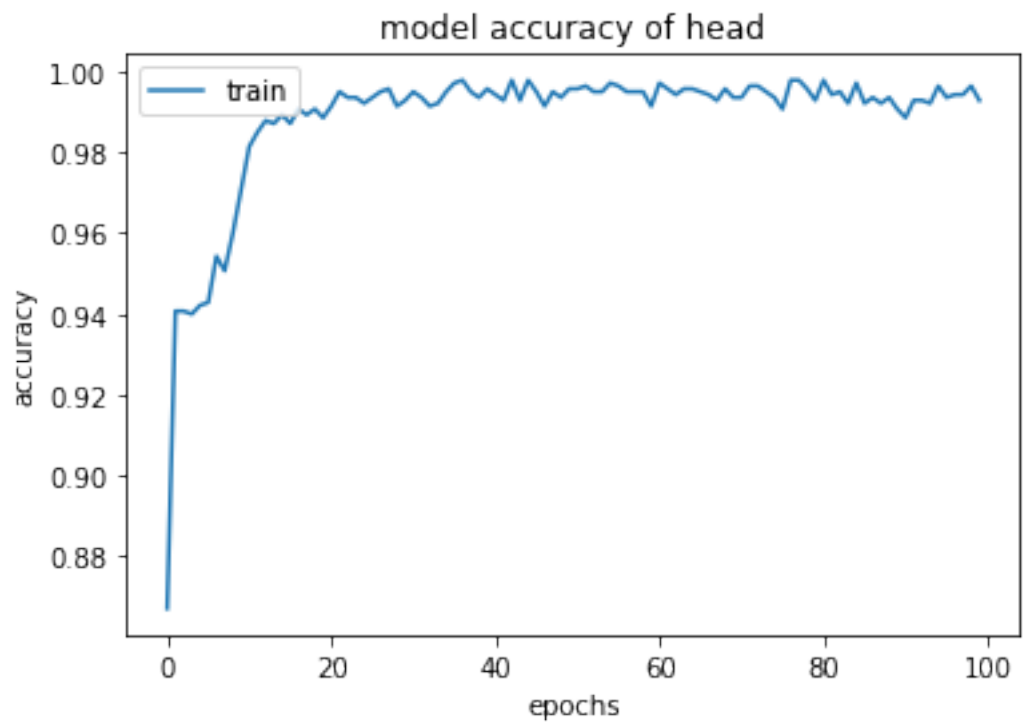


Fig 22

## 4.3 Loss

### Left hip

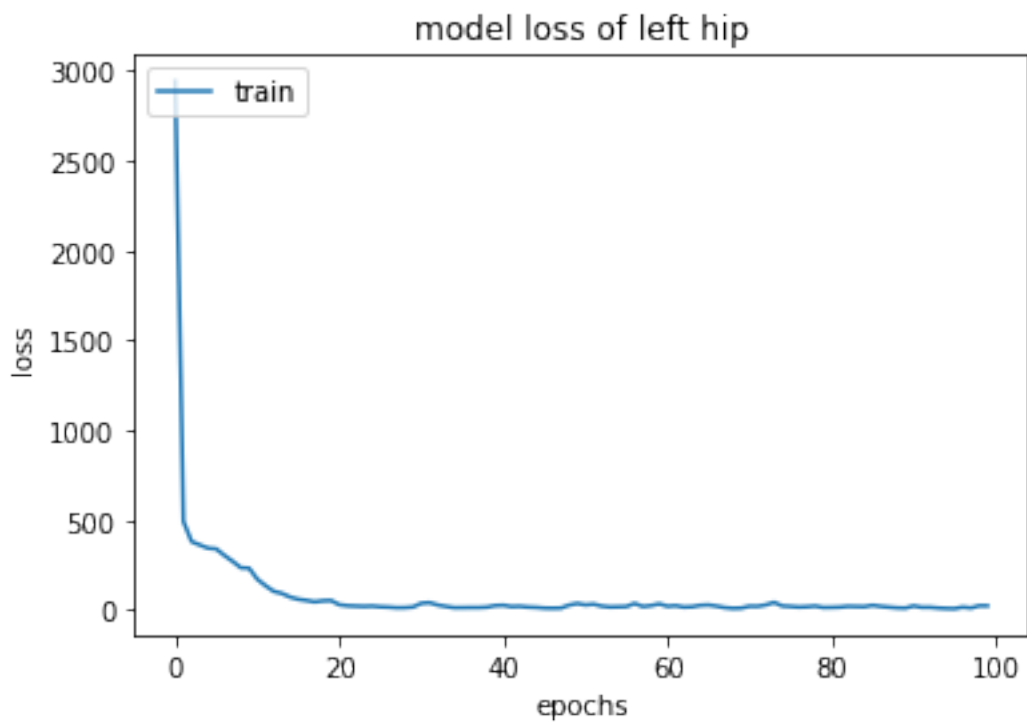


Fig 23

### Left Knee

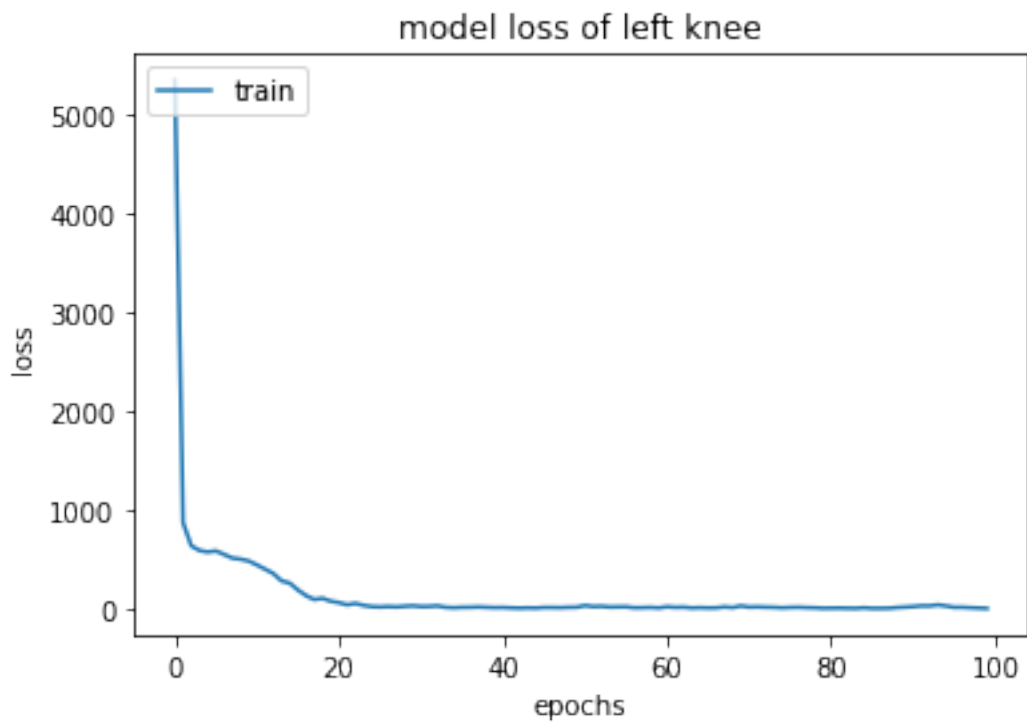


Fig 24

**Left ankle**




Fig 25

**Left shoulder**




Fig 26

**Left elbow**




Fig 27

**Left wrist**

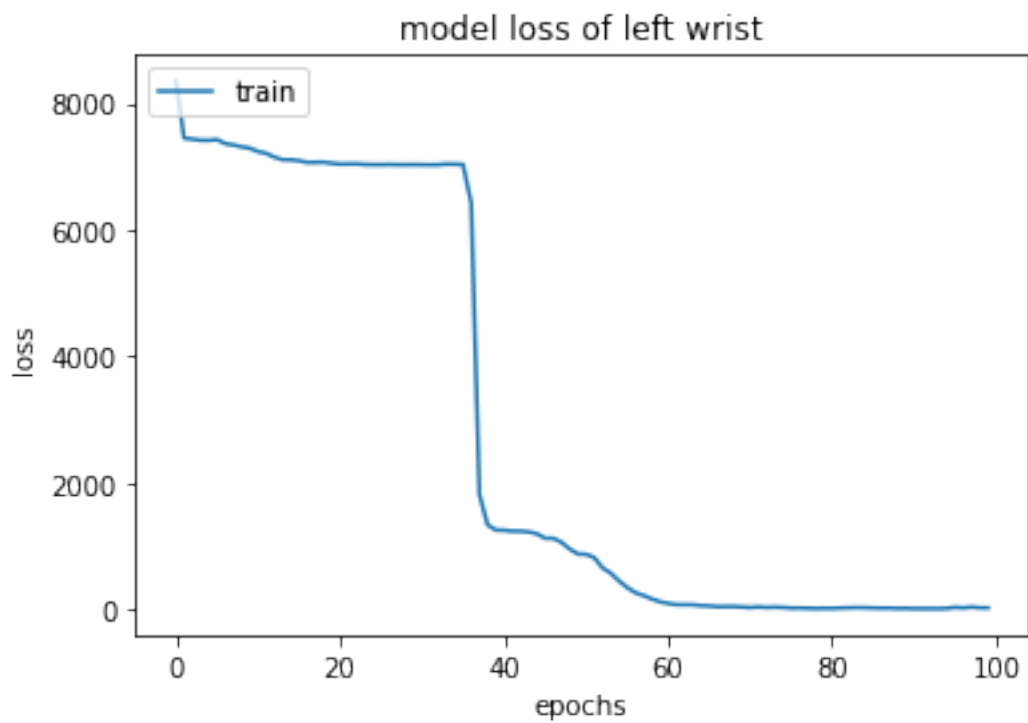


Fig 28



### Right Wrist

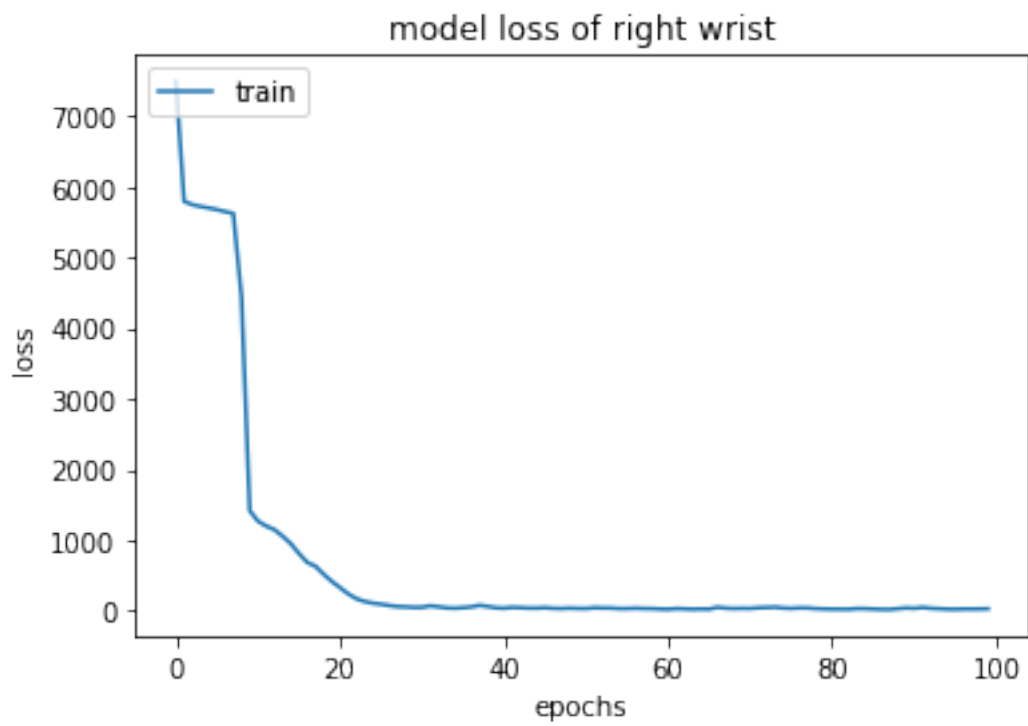


Fig 29

### Right elbow

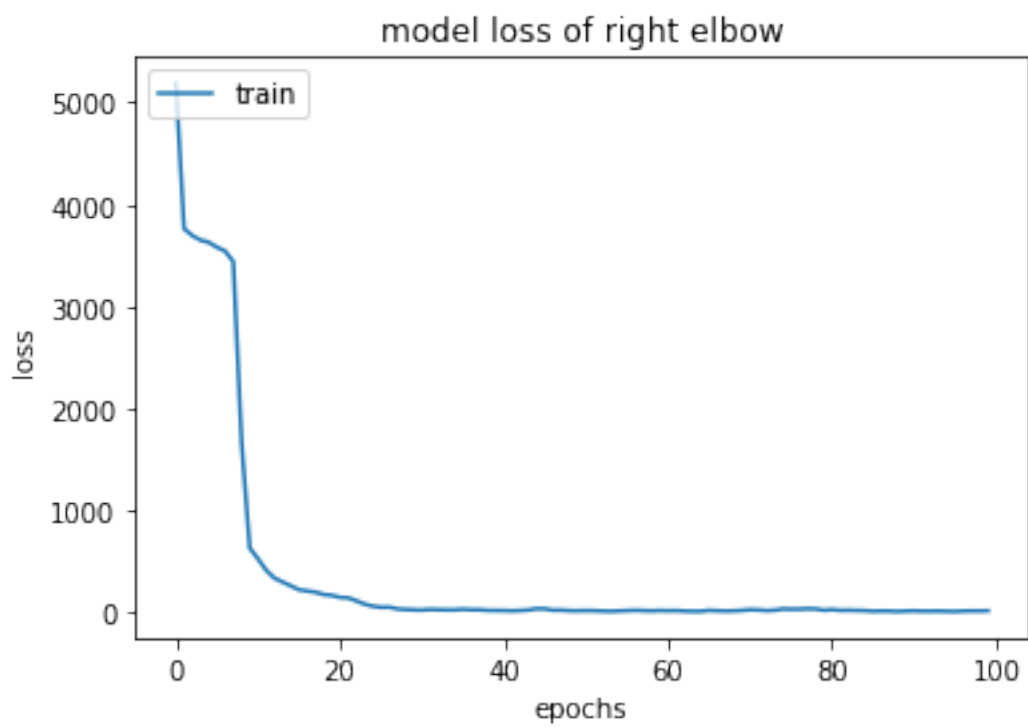


Fig 30

### Right shoulder

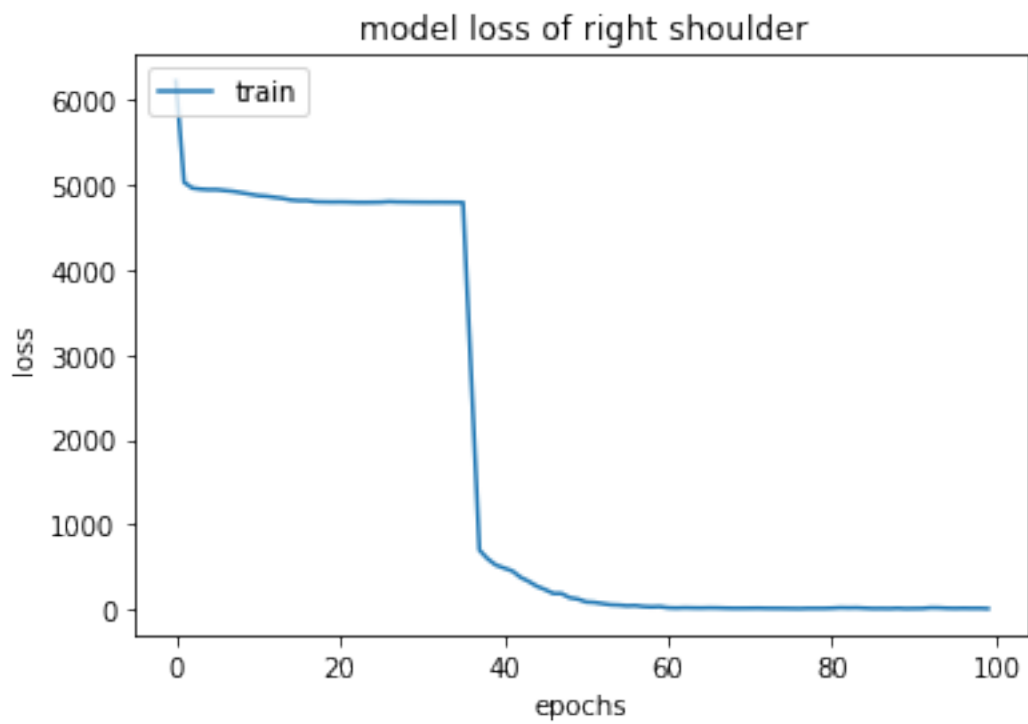


Fig 31

### Right Knee




Fig 32

### Right ankle

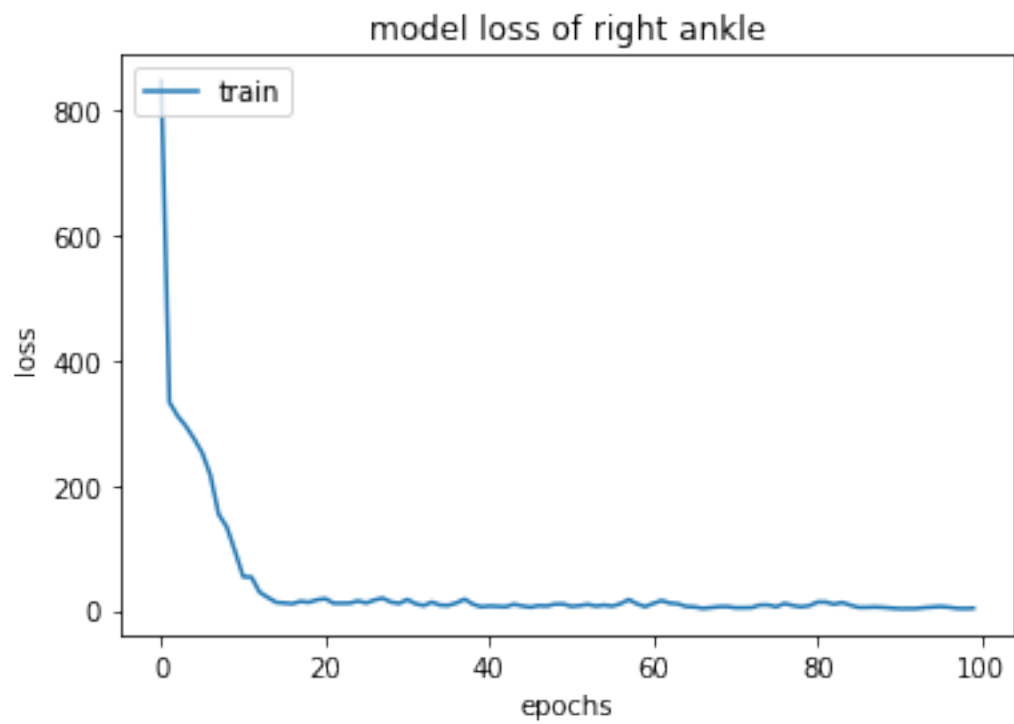


Fig 33

### Right hip




Fig 34

**Neck**

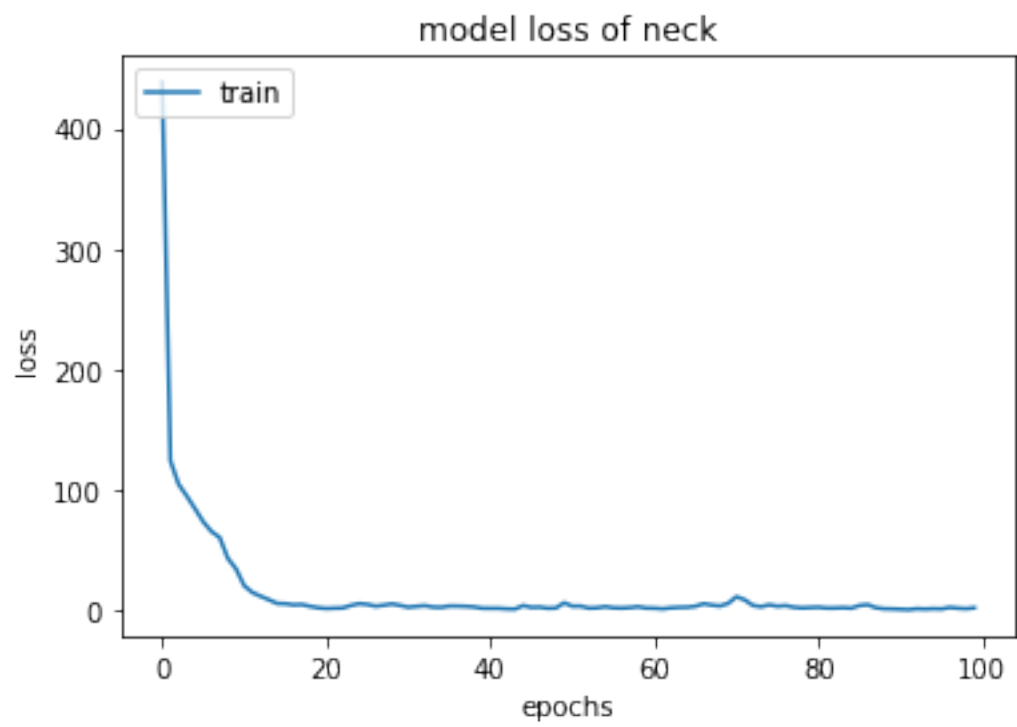


Fig 35

## 4.4 Comparison with other CNN architectures

### 4.4.1 Accuracy of the DeepPose model

#### Final Training and Testing accuracies

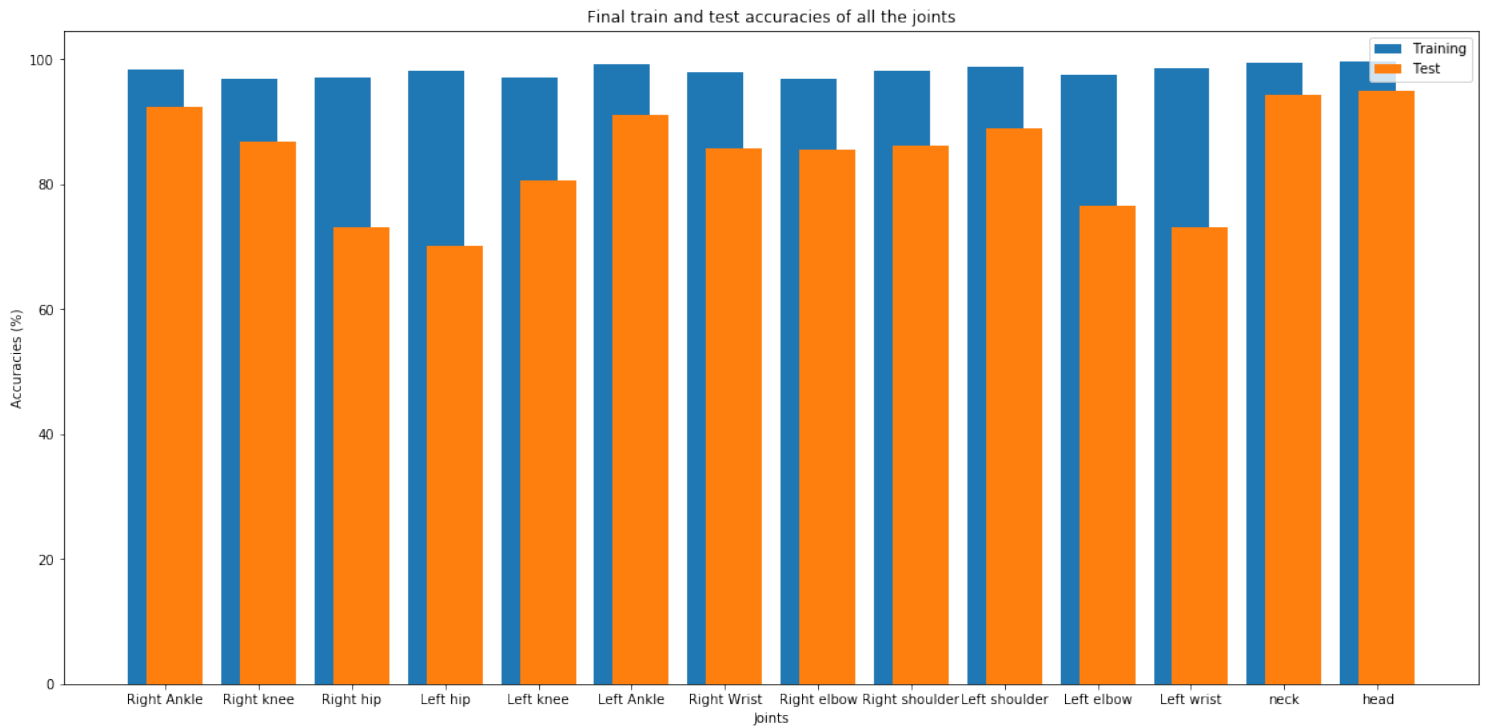


Fig 36

### 4.4.2 Accuracy of the ResNet34 model

#### Final Training and Testing accuracies




Fig 37

### 4.4.3 Accuracy of the AlexNet model

#### Final Training and Testing accuracies




Fig 38

# Chapter 5

## Future Objectives

Following are the future objectives for our project.

### **5.1 Build and real time motion capture application**

After we achieve enough accuracy on the model we plan to build a application that will be able to record real time motion capture data from a simple smartphone camera and this data can be used with many ways like transferring that data to a rigged character to make that virtual character to mimic the movements by animating that character according to the data. It can also be used to create full body AR filters and Controlling cybernetic/robotic arms with the captured data.

# Chapter 6

## Conclusion

This project helped us understand how pose estimation works and how to implement it using deep learning methods. We learnt how to implement it using the different approaches mentioned in the earlier research papers on this topic. Through these implementations, we were also able to come to the conclusion on which method performs the best. Through our own interpretation of these models, we were also able to change certain parameters to improve the performance of the model. We were also able to compare our own model with other existing CNN models, namely, ResNet and AlexNet. We came to the conclusion that different models perform differently based on the joint involved. For better training data, we also created a data annotation tool to get the same.



# References

- [1] Tomas Pfister, James Charles, Andrew Zisserman, *Flowing ConvNets for Human Pose Estimation in Videos*, 2015.
- [2] Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton, *ImageNet Classification with Deep Convolutional Neural Networks*.
- [3] Alexander Toshev, Christian Szegedy, *DeepPose: Human Pose Estimation via Deep Neural Networks*, 2014.
- [4] Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun *Deep Residual Learning for Image Recognition*